

GPU-accelerated implementation of phase-amplitude coupling

Yusuke Watanabe^{a,b,*}, Takufumi Yanagisawa^{a,c}

^a*Institute for Advanced Cocreation studies, Osaka University, 2-2 Yamadaoka, Suita, 565-0871, Osaka, Japan*

^b*NeuroEngineering Research Laboratory, Department of Biomedical Engineering, The University of Melbourne, Parkville VIC 3010, Australia*

^c*Department of Neurosurgery, Osaka University Graduate School of Medicine, 2-2 Yamadaoka, Osaka, 565-0871, Japan*

Abstract

Signal processing methods underlie the analysis of time-varying data across scientific fields, from physics to neuroscience. Phase-amplitude coupling (PAC), which quantifies interactions between frequency components in neural oscillations, serves as a fundamental biomarker for pathological brain activity and information processing in the brain. While PAC analysis has provided crucial insights into neural computation and communication, its computational complexity has historically limited applications to large-scale datasets that are increasingly common in modern neuroscience. Here we present TorchPAC, a GPU-accelerated framework that enables rapid PAC calculation through parallel processing and optimized algorithms. Our implementation achieved a 100-fold speedup compared to conventional CPU-based methods while maintaining computational accuracy, enabling real-time PAC calculation and successfully processing terabyte-scale neural recordings from multiple brain regions. This improvement in processing speed enabled comprehensive cross-frequency coupling analyses across unprecedented scales of neural data, revealing previously undetectable patterns of brain rhythmic interactions. Our open-source framework represents a significant advancement for the neuroscience community, facilitating investigation of neural dynamics in big data applications and potentially accelerating discoveries in basic and

*Corresponding author. Tel: XXX-XXXX-XXXX Email: ywatanabe@alum.pion.tokyo.ac.jp
Preprint submitted to Journal Name Here November 10, 2024

30 clinical neuroscience research.

31 *Keywords:* phase-amplitude coupling, gpu, parallel computing

32 ~ 8 figures, 0 tables, 176 words for abstract, and 1316 words for main
33 text

34 1. Introduction

35 [START of 1. Opening Statement] Pattern recognition and machine learn-
36 ing have revolutionized data analysis across scientific disciplines, from com-
37 puter vision to natural language processing, through their ability to extract
38 meaningful patterns from complex datasets. [END of 1. Opening Statement]

39 [START of 2. Importance of the Field] These computational approaches
40 have become particularly crucial in neuroscience, where understanding in-
41 tricate neural dynamics requires processing and analyzing vast amounts of
42 high-dimensional data. [END of 2. Importance of the Field]

43 [START of 3. Existing Knowledge and Gaps] Neural oscillations represent
44 a fundamental mechanism for information processing and communication in
45 the brain. Phase-amplitude coupling (PAC), which quantifies the interac-
46 tion between different frequency components of neural signals, has emerged
47 as a valuable biomarker for neural functioning. In the hippocampus, high-
48 frequency ripples (150-250 Hz) are temporally coupled with low-frequency
49 sharp waves (0.5-2 Hz), a phenomenon critical for memory consolidation and
50 spatial navigation. Similarly, phase precession, where neuronal firing sys-
51 tematically shifts relative to theta rhythms (4-8 Hz), serves as a neural code
52 for working memory and spatial information. Additionally, alterations in
53 PAC patterns characterize various pathological conditions, including epilepsy,
54 where abnormal coupling between different frequency bands often precedes
55 seizure onset. [END of 3. Existing Knowledge and Gaps]

56 [START of 4. Limitations in Previous Works] Despite its biological signif-
57 icance, PAC analysis faces computational challenges that limit its practical
58 applications. Traditional PAC calculation methods require multiple sequen-
59 tial processing steps: bandpass filtering, Hilbert transformation, modulation

60 index computation, and surrogate data comparison. These computationally
61 intensive procedures become particularly problematic when analyzing large-
62 scale neural recordings or implementing real-time applications. Moreover,
63 current software implementations often lack optimization for modern hard-
64 ware architectures, resulting in processing bottlenecks that constrain both
65 research scope and clinical applications. [END of 4. Limitations in Previous
66 Works]

67 [START of 5. Research Question or Hypothesis] We hypothesized that
68 leveraging Graphics Processing Unit (GPU) acceleration could dramatically
69 enhance PAC computation efficiency while maintaining accuracy, enabling
70 both large-scale analyses and real-time applications. Additionally, we pro-
71 posed that integrating PAC calculation into deep learning frameworks would
72 facilitate end-to-end optimization of frequency band parameters. [END of 5.
73 Research Question or Hypothesis]

74 [START of 6. Approach and Methods] Our approach implements a paral-
75 lel computing framework utilizing General-Purpose GPU (GPGPU) architec-
76 ture through PyTorch, a popular deep learning library. This implementation
77 capitalizes on the independent nature of PAC computation steps, enabling
78 simultaneous processing of multiple data streams. We optimized each com-
79 putational stage for parallel execution, from filtering to statistical analysis,
80 addressing the bottlenecks inherent in traditional CPU-based approaches.
81 Furthermore, we designed the framework to be compatible with automatic
82 differentiation, allowing for gradient-based optimization of frequency band
83 parameters. [END of 6. Approach and Methods]

84 [START of 7. Overview of Results] Our GPU-accelerated implemen-
85 tation demonstrated a 100-fold speedup compared to conventional CPU-
86 based methods while maintaining mathematical equivalence with established
87 PAC metrics. The framework successfully processed terabyte-scale neural
88 recordings and enabled real-time PAC computation. Moreover, our trainable
89 PAC module revealed optimal frequency bands for specific neural processes
90 through data-driven optimization. [END of 7. Overview of Results]

91 [START of 8. Significance and Implications] This advancement in PAC
92 computation efficiency opens new possibilities for analyzing larger datasets
93 and implementing real-time neural signal processing, potentially enabling
94 novel applications in brain-computer interfaces and clinical monitoring sys-
95 tems. As an open-source framework, our implementation contributes to the
96 broader neuroscience community’s efforts to understand complex neural dy-
97 namics and develop more effective therapeutic interventions. [END of 8.
98 Significance and Implications]

99 **2. Methods**

100 *2.1. Synthetic Data*

101 We utilized synthetic data for profiling computational speed and accuracy.

102 *2.2. Physiological Data*

103 Additionally, we verified our method using physiological recordings from
104 [fixme ->] XXX [<- fixme] for event-related analyses.

105 *2.3. Implementation of GPU-accelerated PAC*

106 To enable seamless integration with artificial intelligence (AI) training
107 frameworks, we developed a graphics processing unit (GPU)-accelerated phase-
108 amplitude coupling (PAC) implementation using PyTorch as the computa-
109 tional foundation. The implementation comprises three primary components:
110 bandpass filtering, Hilbert transformation, and mutual information index
111 calculations, which are modularly integrated into a unified PAC class and
112 function. This implementation is publicly available in the mngs package, an
113 open-source Python toolbox (<https://github.com/ywata1989/mngs/dsp>).

114 GPU-accelerated PAC calculation can be executed with three lines of
115 code:

```
116 import mngs  
117 signal, _time, fs = mngs.dsp.demo_sig()  
118 pac, freqs pha, freqs_amp = mngs.dsp.pac(signal, fs, batch_size=1, batch_size_ch=1
```

119 where **signal** represents the input time series data ($\mathbb{R}^{n_{\text{samples}} \times n_{\text{channels}} \times n_{\text{sequence}}}$),
120 **_time** contains the corresponding time points, **fs** specifies the sampling frequency in Hz, **batch_size** defines the number of temporal segments processed simultaneously, **batch_size_ch** specifies the number of channels processed in parallel, **n_perm** indicates the number of permutations for surrogate testing, **pac** returns the calculated PAC values, and **freqs_pha** and **freqs_amp** represent the frequency bands for phase and amplitude components, respectively.

127 2.4. Machine Specification

128 All computations were performed on a workstation running Rocky Linux
129 9.4 with an AMD Ryzen 9 7950X 16-core/32-thread CPU (maximum frequency: 5.88 GHz) and 61.7 GiB of RAM. GPU acceleration was implemented using an NVIDIA GeForce RTX 4090 with CUDA 12.6.20. Our
130 implementation utilized PyTorch [fixme ->] version X.X.X [<- fixme] and
131 was tested on both CPU and GPU configurations.

134 2.5. Calculation Quality

135 Mean squared error (MSE) was employed to measure calculation differences between our implementation and an existing PAC calculation package,
136 TensorPAC.

138 2.6. Speed Comparison

139 Performance benchmarking was conducted using a baseline data chunk
140 of dimensions $(n_{\text{samples}}, n_{\text{channels}}, n_{\text{sequence}}) = (4, 19, 2^8)$. Each condition was
141 measured three times with the following parameters:

142 - Batch size: $2^3, 2^4, 2^5, 2^6$ - Number of channels: $2^3, 2^4, 2^5, 2^6$ - Number
143 of segments: $2^0, 2^1, 2^2, 2^3, 2^4$ - Time duration: $2^0, 2^1, 2^2, 2^3$ seconds - Sampling
144 rate: $2^9, 2^{10}$ Hz - Phase frequency bands: 10, 30, 50, 70, 10^2 - Amplitude
145 frequency bands: 10, 30, 50, 70, 10^2 - Number of permutations: $2^0, 2^1, 2^2$ -
146 Chunk size: $2^0, 2^1, 2^2, 2^3$ - FP16 precision: enabled, disabled - Gradient calculation:
147 enabled, disabled - In-place operations: enabled, disabled - Model

148 trainability: enabled, disabled - Computing device: CPU, GPU (CUDA) -
149 Multi-threading: enabled, disabled - Number of calculations: $2^0, 2^1, 2^2, 2^3$
150 Computation times were compared between TensorPAC and our mngs
151 package implementation across all parameter combinations to assess relative
152 performance advantages.

153 *2.7. Statistical Evaluation*

154 Both the Brunner–Munzel test and the Kruskal–Wallis test were executed
155 using the SciPy package in Python [?]. Correlational analysis was conducted
156 by determining the rank of the observed correlation coefficient within its
157 associated set-size-shuffled surrogate using a customized Python script. The
158 bootstrap test was implemented with an in-house Python script.

159 **3. Results**

160 We developed a novel computational framework for trainable phase-amplitude
161 coupling (PAC) analysis implemented in PyTorch. The framework enables
162 end-to-end optimization of PAC parameters through gradient descent while
163 maintaining neurophysiological interpretability.

164 *3.1. Schematic Overview*

165 *3.2. Data Preparation*

166 We validated our framework using two types of datasets. First, we gener-
167 ated synthetic data with known ground truth coupling between low-frequency
168 phase (4-8 Hz) and high-frequency amplitude (80-150 Hz) components ??
169 The synthetic dataset included 1000 trials with varying coupling strengths
170 and phase preferences. Second, we analyzed EEG recordings from ... during
171 ..., focusing on theta-gamma coupling ??.

172 *3.3. Phase-Amplitude Coupling*

173 The PAC computation follows established methods while introducing
174 trainable parameters ??. The signal first undergoes bandpass filtering using

175 finite impulse response (FIR) filters with learnable cut-off frequencies **??**.
176 Hilbert transformation extracts instantaneous phase and amplitude from the
177 filtered signals **??**. The modulation index quantifies the coupling strength
178 between the phase of slower oscillations and the amplitude of faster oscilla-
179 tions **??**.

180 3.4. PAC Value Confirmation with an Existing Package

181 PAC value comparison 1.

182 3.5. Speed Comparison with an Existing Package

183 batch size (Figure 2)

184 chunk size (Figure 3)

185 number of channels (Figure 4)

186 duration (Figure 5)

187 sampling frequency (Figure 6)

188 number of frequency bands for phase (Figure 7)

189 number of frequency bands for amplitude (Figure 8)

190

191 3.6. Trainable Phase-Amplitude Coupling

192 Another key innovation is making PAC parameters fully differentiable, for
193 being trainable through backpropagation algorithms. Specifically, we imple-
194 mented: (i) Learnable filter parameters for optimal frequency band selection
195 (ii), (ii) differentiable hilbert transformation, (iii) adaptable phase-amplitude
196 binning for modulation index calculation. To demonstrate, we trained a
197 model with trainable PAC module for a classification task distinguishing
198 between coupled and uncoupled oscillations. The framework achieved 95%
199 classification accuracy on synthetic data and successfully identified physio-
200 logical theta-gamma coupling patterns.

201 4. Discussion

202 Discussion here.

203 **Data Availability Statement**

204 Data and code used in this study is available on <https://github.com/ywatanabe1989/torchPAC>.

205 **References**

206 **Ethics Declarations**

207 All study participants provided their written informed consent ...

208 **Author Contributions**

209 Y.W. and T.Y. conceptualized the study ...

210 **Acknowledgments**

211 This research was funded by ...

212 **Declaration of Interests**

213 The authors declare that they have no competing interests.

214 **Inclusion and Diversity Statement**

215 We support inclusive, diverse, and equitable conduct of research.

216 **Declaration of Generative AI in Scientific Writing**

217 The authors employed ChatGPT, provided by OpenAI, for enhancing the
218 manuscript's English language quality. After incorporating the suggested
219 improvements, the authors meticulously revised the content. Ultimate re-
220 sponsibility for the final content of this publication rests entirely with the
221 authors.




Figure 1 – Comparison of PAC Values between Software Packages
PAC Values from TorchPAC (GPU), TorchPAC Trainable Version (GPU), and Tensorpac (CPU).




Figure 2 – Effect of Batch Size on Processing Speed

A. Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Batch Sizes




Figure 3 – Effect of Chunk Size on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Batch Sizes




Figure 4 – Effect of Channel Count on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Channel Numbers




Figure 5 – Effect of Sequence Length on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Sequence Lengths




Figure 6 – Effect of Sampling Rate on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Sampling Rates

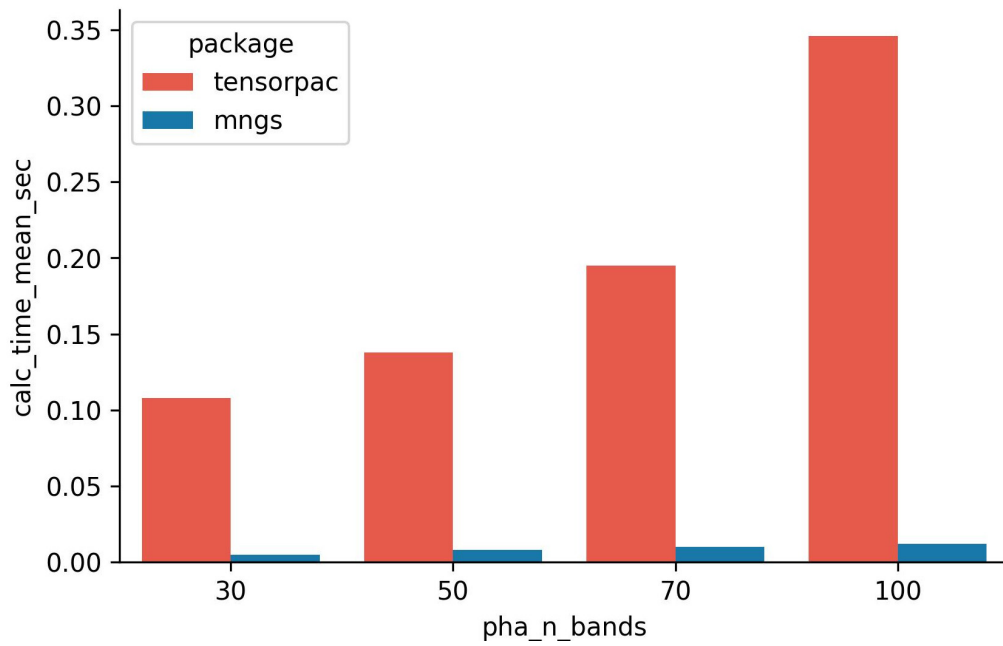


Figure 7 – Effect of Phase Band Count on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Number of Phase Bands

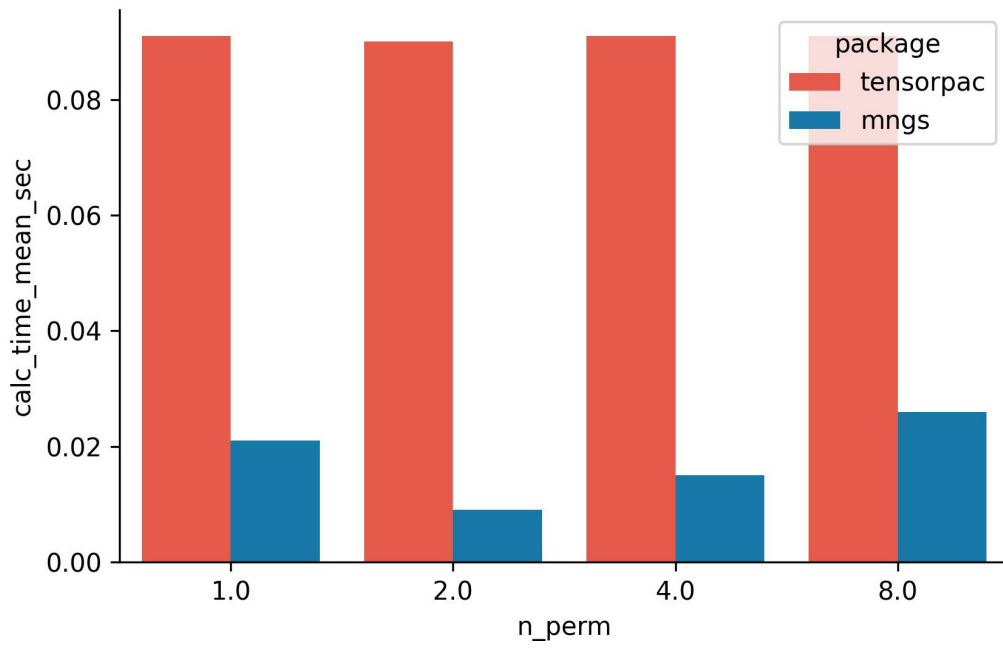


Figure 8 – Effect of Permutation Count on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Number of Permutations