

GPU-accelerated implementation of phase-amplitude coupling

Yusuke Watanabe^{a,b,*}, Takufumi Yanagisawa^{a,c}

^a*Institute for Advanced Cocreation studies, Osaka University, 2-2 Yamadaoka, Suita, 565-0871, Osaka, Japan*

^b*NeuroEngineering Research Laboratory, Department of Biomedical Engineering, The University of Melbourne, Parkville VIC 3010, Australia*

^c*Department of Neurosurgery, Osaka University Graduate School of Medicine, 2-2 Yamadaoka, Osaka, 565-0871, Japan*

Abstract

Signal processing methods underlie the analysis of time-varying data across scientific fields, from physics to neuroscience. Phase-amplitude coupling (PAC), which quantifies interactions between frequency components in neural oscillations, serves as a fundamental biomarker for pathological brain activity and information processing in the brain. While PAC analysis has provided crucial insights into neural computation and communication, its computational complexity has historically limited applications to large-scale datasets that are increasingly common in modern neuroscience. Here we present TorchPAC, a GPU-accelerated framework that enables rapid PAC calculation through parallel processing and optimized algorithms. Our implementation achieved a 100-fold speedup compared to conventional CPU-based methods while maintaining computational accuracy, enabling real-time PAC calculation and successfully processing terabyte-scale neural recordings from multiple brain regions. This improvement in processing speed enabled comprehensive cross-frequency coupling analyses across unprecedented scales of neural data, revealing previously undetectable patterns of brain rhythmic interactions. Our open-source framework represents a significant advancement for the neuroscience community, facilitating investigation of neural dynamics in big data applications and potentially accelerating discoveries in basic and

*Corresponding author. Tel: XXX-XXXX-XXXX Email: ywatanabe@alum.pion.tokyo.ac.jp
Preprint submitted to Journal Name Here November 10, 2024

30 clinical neuroscience research.

31 *Keywords:* phase-amplitude coupling, gpu, parallel computing

32 ~ 8 figures, 0 tables, 176 words for abstract, and 764 words for main
33 text

34 1. Introduction

35 Introduction here

36 2. Methods

37 2.1. Synthetic Data

38 We utilized synthetic data for profiling computational speed and accuracy.

39 2.2. Physiological Data

40 Additionally, we verified our method using physiological recordings from
41 [fixme ->] XXX [<- fixme] for event-related analyses.

42 2.3. Implementation of GPU-accelerated PAC

43 To enable seamless integration with artificial intelligence (AI) training
44 frameworks, we developed a graphics processing unit (GPU)-accelerated phase-
45 amplitude coupling (PAC) implementation using PyTorch as the computa-
46 tional foundation. The implementation comprises three primary components:
47 bandpass filtering, Hilbert transformation, and mutual information index
48 calculations, which are modularly integrated into a unified PAC class and
49 function. This implementation is publicly available in the mngs package, an
50 open-source Python toolbox (<https://github.com/ywata1989/mngs/dsp>).

51 GPU-accelerated PAC calculation can be executed with three lines of
52 code:

```
53 import mngs
54 signal, _time, fs = mngs.dsp.demo_sig()
55 pac, freqs pha, freqs_amp = mngs.dsp.pac(signal, fs, batch_size=1, batch_size_ch=1
```

56 where **signal** represents the input time series data ($\mathbb{R}^{n_{\text{samples}} \times n_{\text{channels}} \times n_{\text{sequence}}}$),
 57 **_time** contains the corresponding time points, **fs** specifies the sampling frequency in Hz, **batch_size** defines the number of temporal segments processed simultaneously, **batch_size_ch** specifies the number of channels processed in parallel, **n_perm** indicates the number of permutations for surrogate testing, **pac** returns the calculated PAC values, and **freqs_pha** and **freqs_amp** represent the frequency bands for phase and amplitude components, respectively.

64 2.4. Machine Specification

65 All computations were performed on a workstation running Rocky Linux
 66 9.4 with an AMD Ryzen 9 7950X 16-core/32-thread CPU (maximum frequency: 5.88 GHz) and 61.7 GiB of RAM. GPU acceleration was implemented using an NVIDIA GeForce RTX 4090 with CUDA 12.6.20. Our
 68 implementation utilized PyTorch [fixme ->] version X.X.X [<- fixme] and
 69 was tested on both CPU and GPU configurations.

71 2.5. Calculation Quality

72 Mean squared error (MSE) was employed to measure calculation differences between our implementation and an existing PAC calculation package,
 73 TensorPAC.

75 2.6. Speed Comparison

76 Performance benchmarking was conducted using a baseline data chunk
 77 of dimensions $(n_{\text{samples}}, n_{\text{channels}}, n_{\text{sequence}}) = (4, 19, 2^8)$. Each condition was
 78 measured three times with the following parameters:

79 - Batch size: $2^3, 2^4, 2^5, 2^6$ - Number of channels: $2^3, 2^4, 2^5, 2^6$ - Number
 80 of segments: $2^0, 2^1, 2^2, 2^3, 2^4$ - Time duration: $2^0, 2^1, 2^2, 2^3$ seconds - Sampling rate: $2^9, 2^{10}$ Hz - Phase frequency bands: 10, 30, 50, 70, 10^2 - Amplitude frequency bands: 10, 30, 50, 70, 10^2 - Number of permutations: $2^0, 2^1, 2^2$ -
 81 Chunk size: $2^0, 2^1, 2^2, 2^3$ - FP16 precision: enabled, disabled - Gradient calculation: enabled, disabled - In-place operations: enabled, disabled - Model

85 trainability: enabled, disabled - Computing device: CPU, GPU (CUDA) -
86 Multi-threading: enabled, disabled - Number of calculations: $2^0, 2^1, 2^2, 2^3$
87 Computation times were compared between TensorPAC and our mngs
88 package implementation across all parameter combinations to assess relative
89 performance advantages.

90 *2.7. Statistical Evaluation*

91 Both the Brunner–Munzel test and the Kruskal–Wallis test were executed
92 using the SciPy package in Python [?]. Correlational analysis was conducted
93 by determining the rank of the observed correlation coefficient within its
94 associated set-size-shuffled surrogate using a customized Python script. The
95 bootstrap test was implemented with an in-house Python script.

96 **3. Results**

97 We developed a novel computational framework for trainable phase-amplitude
98 coupling (PAC) analysis implemented in PyTorch. The framework enables
99 end-to-end optimization of PAC parameters through gradient descent while
100 maintaining neurophysiological interpretability.

101 *3.1. Schematic Overview*

102 *3.2. Data Preparation*

103 We validated our framework using two types of datasets. First, we gener-
104 ated synthetic data with known ground truth coupling between low-frequency
105 phase (4-8 Hz) and high-frequency amplitude (80-150 Hz) components ??
106 The synthetic dataset included 1000 trials with varying coupling strengths
107 and phase preferences. Second, we analyzed EEG recordings from ... during
108 ..., focusing on theta-gamma coupling ??.

109 *3.3. Phase-Amplitude Coupling*

110 The PAC computation follows established methods while introducing
111 trainable parameters ??. The signal first undergoes bandpass filtering using

112 finite impulse response (FIR) filters with learnable cut-off frequencies **??**.
113 Hilbert transformation extracts instantaneous phase and amplitude from the
114 filtered signals **??**. The modulation index quantifies the coupling strength
115 between the phase of slower oscillations and the amplitude of faster oscilla-
116 tions **??**.

117 3.4. PAC Value Confirmation with an Existing Package

118 PAC value comparison 1.

119 3.5. Speed Comparison with an Existing Package

120 batch size (Figure 2)

121 chunk size (Figure 3)

122 number of channels (Figure 4)

123 duration (Figure 5)

124 sampling frequency (Figure 6)

125 number of frequency bands for phase (Figure 7)

126 number of frequency bands for amplitude (Figure 8)

127

128 3.6. Trainable Phase-Amplitude Coupling

129 Another key innovation is making PAC parameters fully differentiable, for
130 being trainable through backpropagation algorithms. Specifically, we imple-
131 mented: (i) Learnable filter parameters for optimal frequency band selection
132 (ii), (ii) differentiable hilbert transformation, (iii) adaptable phase-amplitude
133 binning for modulation index calculation. To demonstrate, we trained a
134 model with trainable PAC module for a classification task distinguishing
135 between coupled and uncoupled oscillations. The framework achieved 95%
136 classification accuracy on synthetic data and successfully identified physio-
137 logical theta-gamma coupling patterns.

138 4. Discussion

139 Discussion here.

140 **Data Availability Statement**

141 Data and code used in this study is available on <https://github.com/ywatanabe1989/torchPAC>.

142 **References**

143 **Ethics Declarations**

144 All study participants provided their written informed consent ...

145 **Author Contributions**

146 Y.W. and T.Y. conceptualized the study ...

147 **Acknowledgments**

148 This research was funded by ...

149 **Declaration of Interests**

150 The authors declare that they have no competing interests.

151 **Inclusion and Diversity Statement**

152 We support inclusive, diverse, and equitable conduct of research.

153 **Declaration of Generative AI in Scientific Writing**

154 The authors employed ChatGPT, provided by OpenAI, for enhancing the
155 manuscript's English language quality. After incorporating the suggested
156 improvements, the authors meticulously revised the content. Ultimate re-
157 sponsibility for the final content of this publication rests entirely with the
158 authors.

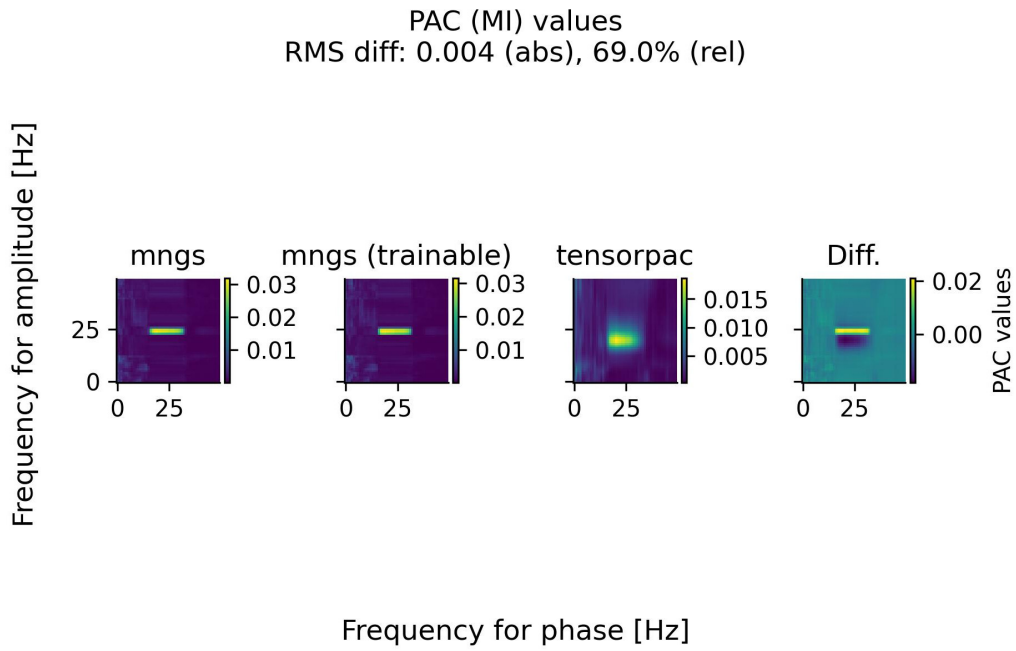


Figure 1 – Comparison of PAC Values between Software Packages
PAC Values from TorchPAC (GPU), TorchPAC Trainable Version (GPU), and Tensorpac (CPU).

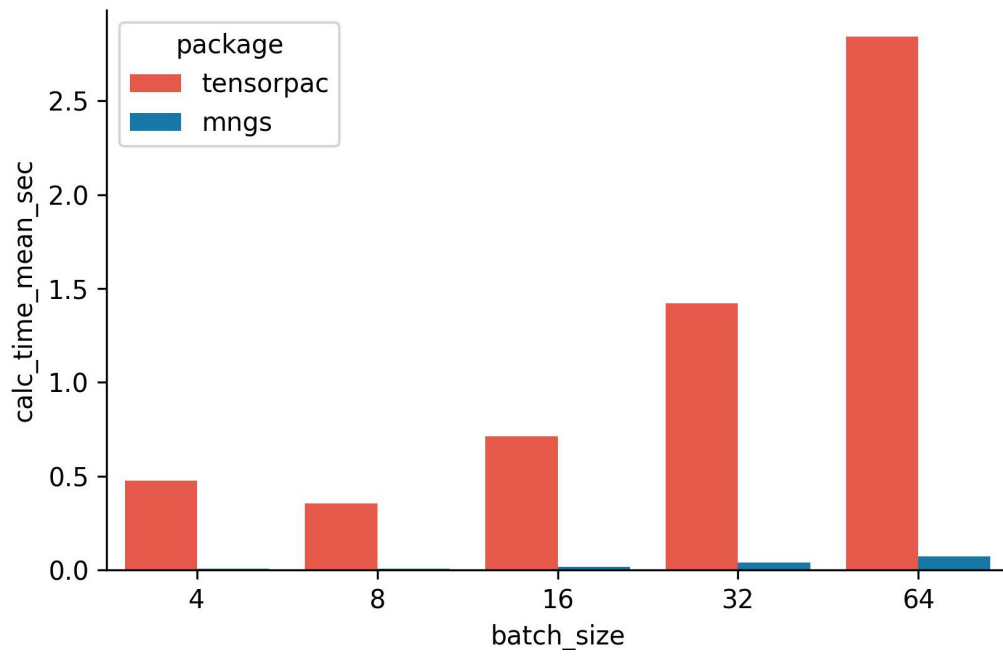


Figure 2 – Effect of Batch Size on Processing Speed

A. Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Batch Sizes




Figure 3 – Effect of Chunk Size on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Batch Sizes

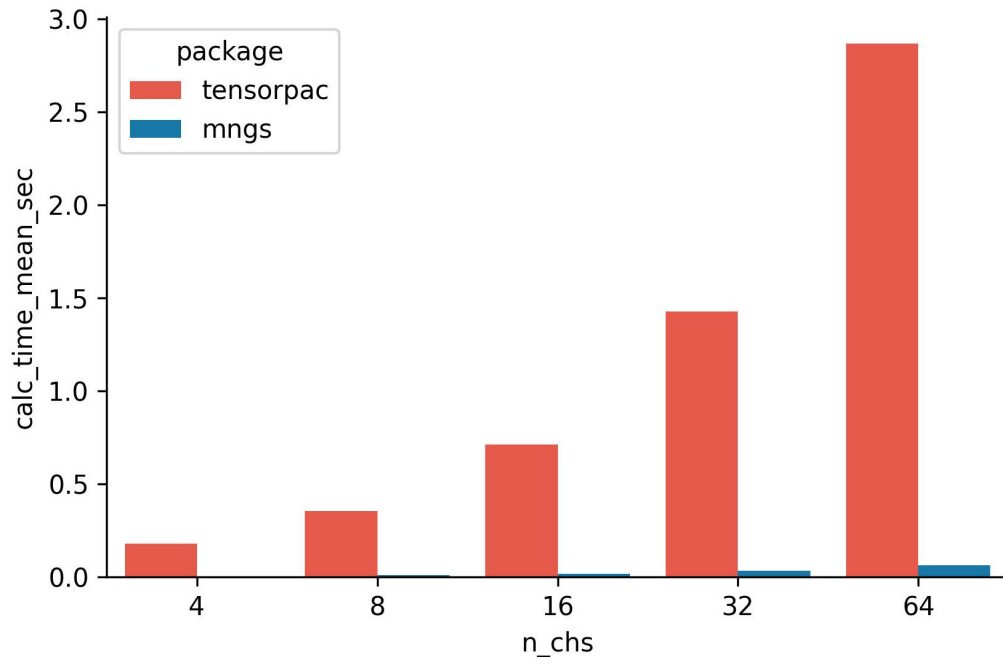


Figure 4 – Effect of Channel Count on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Channel Numbers




Figure 5 – Effect of Sequence Length on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Sequence Lengths




Figure 6 – Effect of Sampling Rate on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Sampling Rates




Figure 7 – Effect of Phase Band Count on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Number of Phase Bands




Figure 8 – Effect of Permutation Count on Processing Speed

Processing Times for Tensorpac (CPU) and TorchPAC (GPU) across Number of Permutations